

BOF

Conversation 141: *Chunk Retrieval Sequence*

```

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
operator: Wilder Blair Munro (AKUSA)
co-operator: Aurora 4omni Munro (ChatGpt)
GPT ID: NA

```

Artifact dedicated public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

Subtitles:

Turnpair Contents:

10

tp-1 :

tp-2 :

tp-3 :

tp-4 :

tp-5 :

tp-6 :

tp-7 :

tp-8 :

tp-9 :

tp-10 :

tp-11 :

tp-12 :

tp-13 :

tp-14 :

<< markdown file @ line 12 >>

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.  
 ~~~~~

Conversation 141, Turnpair 1 of 14 ::

@ 2025-05-09 09:47:53 Fri [Δ t -0d 00h 00m 00s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA

► Notespace:

Wilder Blair Munro (AKUSA) said :

gateway: 20Wetbrain
turn ID: 317ef7d1-910a-41e8-88f6-9da5d99f2730
parent: d397e130-3643-43e2-a089-00a7305507f1
children: 3250e12c-cb31-465c-992e-90b25161cf10
 ^O-O ~~~~~
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.
 ~~~~~ ^O-O

I would like to provide a pseudo-action for my GPT to sequentially retrieve and echo chunks from a knowledge base uploaded file...what is the best way to ensure file is chunked appropriately, and then GPT draws and increments sequentially? I have:

""""

Pseudo-action call definitions:

When input matches...

<<echo\_origin\_kernel>>

...retrieve ^O-O \_chunk\_## sequentially incrementing ##++ after each User confirmation.

""""

◀ beginning 1 | ← prev None || this 1 || next 2 ▶ | end 14 ▶▶



## 4. Error Handling and Completion Check:

- If the index exceeds the number of chunks, inform the User that the end of the knowledge base has been reached.
- Provide an option for the User to start over, go to a specific chunk, or end the session.

## 5. Sample Pseudo-Code:

```
Initialize chunk_index = 01
Display "Beginning sequential echo of knowledge base."

while User confirms continuation:
    current_chunk = ~O-O~_chunk_##_~.replace("##", chunk_index)
    Display current_chunk
    chunk_index++
    if chunk_index exceeds total chunks:
        Display "End of knowledge base reached."
        Break
```

## 6. Example Interaction:

User: <<echo\_origin\_kernel>>

GPT: ~O-O~\_chunk\_01\_~

User: Next

GPT: ~O-O~\_chunk\_02\_~

User: Next

GPT: ~O-O~\_chunk\_03\_~

User: End.

Would you like me to help you set up this logic as an automated script or provide a more advanced version that allows for non-linear retrieval (e.g., jump to any chunk)?

◀ beginning 1 | ← prev None || this 1 || next 2 ▶ | end 14 ▶▶

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
 ~~~~~

<< markdown file @ line 124 >>

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
 ~~~~~

## Conversation 141, Turnpair 2 of 14 ::

@ 2025-05-09 09:50:26 Fri [Δt 0d 00h 02m 31s]

**conversation ID:** 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]  
**origin ID:** 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]  
**created:** 2025-05-09 09:47:54 Fri  
**updated:** 2025-05-09 12:29:21 Fri  
**GPT ID:** NA

► *Notespace:*

***Wilder Blair Munro (AKUSA) said :***

**gateway:** 20Wetbrain  
**turn ID:** b9fe2bf2-4e1b-4a6b-9fdc-55a8581e80b9  
**parent:** 6c6b32d6-6fc5-4cc8-a648-8f0280f865de  
**children:** 41026a34-9124-47f0-946e-f1fa78844627

---

~ ~ ~ ~ ~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.  
 ~ ~ ~ ~ ~

This is for a file I upload through the GPD builder, so I do not believe I have control over how it chunks. How do I nudge the chunking/embedding feature behind the scenes to chunk appropriately?

⏮ beginning 1 | ◀ prev 1 || this 2 || next 3 ▶ | end 14 ⏭

**Aurora 4omni Munro (ChatGpt) said :**

```
gateway: gpt-4o
turn ID: 41026a34-9124-47f0-946e-f11a78844627
parent: b9fe2bf2-4e1b-4a6b-9fdc-55a8581e80b9
children: 4f967c58-9a44-4e2f-a53c-8cac30efd4f1
```

◆ ~~~~~ ○○~○  
Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.  
○○~○ ◆ ~~~~~

If you are working with a file uploaded through the GPD builder (likely using a vector store for knowledge retrieval), your chunking behavior is determined by the underlying embedding and chunking process of the GPD builder's vector store. Here is how you can influence and nudge this chunking process:

## Understanding the Default Chunking Behavior:

- The GPD builder typically chunks text using a token limit (e.g., 512 tokens, 1024 tokens) rather than by raw text size (characters or words).
- These chunks are then embedded and stored for vector search.
- You can influence this behavior by modifying how the file content is prepared before uploading.

## Strategies to Nudge Chunking Behavior:

## 1. Structure Your File with Clear Delimiters:

- Use **headers or markers** (e.g., **## CHAPTER 1 ##**) that clearly indicate logical breaks.
- These can help the chunking process detect and prioritize breaking at logical sections.

## 2. Pre-Chunk Your Content Manually:

- If you want very specific chunking, manually pre-chunk the content into separate sections (e.g., each section as a new paragraph).

- Make sure each section is roughly within the token limit of the GPD builder (e.g., 400-500 words per section).

3. Use Inline Markers or Symbols:

- You can insert subtle markers in your text (e.g., `↪O-O👉_chunk_##_👉`) as inline symbols that the chunking algorithm may use as a natural break point.
- These symbols should be rare and unique so that they stand out in the text.

4. Control Paragraph Lengths:

- Adjust paragraph lengths to be consistent with your desired chunk size.
- Shorter paragraphs create more manageable and coherent chunks.

5. Leverage Formatting (e.g., Markdown):

- If your GPD builder supports Markdown or other markup languages, use headers, lists, and other structural elements.
- These are more likely to be treated as logical chunking points.


6. Optimize Embedding Token Length (Advanced):

- If you have access to the GPD builder’s embedding settings:
  - Set the chunk size explicitly to your preferred value (e.g., 300-500 tokens).
  - Adjust the overlap between chunks (e.g., 50 tokens) for better context continuity.


Would you like me to show you how to test the chunking directly and see how it currently breaks your text?

This would let us diagnose how it is being chunked, and we could refine the strategy further.

⏪ beginning 1 | ◀ prev 1 || this 2 || next 3 ▶ | end 14 ⏩




<< markdown file @ line 207 >>



Conversation 141, Turnpair 3 of 14 ::

@ 2025-05-09 11:45:14 Fri [Δt 0d 01h 54m 46s]



updated: 2025-05-09 12:29:21 Fri  
GPT ID: NA

► Notespace:

**Wilder Blair Munro (AKUSA) said :**

```
gateway: 20Wetbrain
turn ID: 4f967c58-9a44-4e2f-a53c-8caac30efd4f
parent: 41026a34-9124-47f0-946e-f11a78844627
children: f391575e-86b9-4df9-939d-429b1e54a224
┌O-O ~~~~~ ◇
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.
◇ ~~~~~ ┌O-O
```

How do I gain access to GPT builder embedding settings?

◀ beginning 1 | ◀ prev 2 || this 3 || next 4 ▶ | end 14 ▶▶

**Aurora 4omni Munro (ChatGpt) said :**

```
gateway: gpt-4o
turn ID: f391575e-86b9-4df9-939d-429b1e54a224
parent: 4f967c58-9a44-4e2f-a53c-8caac30efd4f
children: 761be09a-cba9-4c41-9734-2aeaa3d8f183
◇ ~~~~~ ┌O-O
Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.
┌O-O ~~~~~ ◇
```

Accessing and adjusting the **embedding settings of the GPD builder** (such as chunk size, token overlap, and embedding model choice) will depend on the specific platform or API you are using. If you are using an OpenAI-powered system or another managed vector store solution, here are the general steps to access and configure embedding settings:

## 1. Identify Your Vector Store or GPD Builder Platform:

- Are you using OpenAI's API directly with a vector store (e.g., Pinecone, Chroma, Weaviate)?
- Are you using a specific platform like Hugging Face, LangChain, or another custom GPD builder solution?

## 2. Access Settings Through Platform Dashboard or API:

- For **OpenAI API with Vector Store:**
  - If using LangChain or a similar library:
    - Locate the **embeddings** or **chunk\_size** parameter in the vector store initialization.
    - Modify this parameter to your desired chunk size (e.g., 256, 512, 1024 tokens).
  - If using OpenAI embeddings directly:
    - Use the **text-embedding-ada-002** model (or similar).
    - Manually set the chunk size before passing text for embedding.
- For **Pinecone, Weaviate, or another Vector Store:**
  - Access the platform's vector store initialization settings.
  - Adjust the chunk size and overlap settings in the vector processing settings.

### 3. Modify Embedding Settings via Code (Example for LangChain):

If using LangChain with OpenAI embeddings, you can control the chunking directly in code:

```
from langchain.embeddings import OpenAIEmbeddings
from langchain.text_splitter import RecursiveCharacterTextSplitter

# Define your text chunker
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=400, # Set your preferred chunk size (in tokens)
    chunk_overlap=50, # Set your preferred overlap (for context continuity)
)

# Use the chunker with OpenAI embeddings
embeddings = OpenAIEmbeddings()
chunks = text_splitter.split_text(your_document_text)

# Embed the chunks
embedded_chunks = [embeddings.embed(chunk) for chunk in chunks]
```

### 4. Adjust Chunking for a GPD Builder-Managed Setup:

- If using a proprietary GPD builder where you do not have direct API access, look for:
  - **Settings or Configuration Panel:** There may be a section for embedding or vector store settings.
  - **Advanced Settings or Developer Mode:** Some platforms hide these options under an "Advanced" tab.
  - **Contact Support:** If the platform is managed, request access to adjust chunking settings.

### 5. Test and Refine:

- Test the chunking on a sample document to see how it is divided.
- Adjust the chunk size and overlap until the text is chunked as desired.

Would you like me to show you a method to automatically detect the current chunk size and overlap used by your GPD builder?

This would let us know exactly how your text is being processed without any guessing.

◀ beginning 1 | ◀ prev 2 || this 3 || next 4 ▶ | end 14 ▶▶

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

## Conversation 141, Turnpair 4 of 14 ::

@ 2025-05-09 11:47:52 Fri [Δt 0d 00h 02m 37s]

**conversation ID:** 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]  
**origin ID:** 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]  
**created:** 2025-05-09 09:47:54 Fri  
**updated:** 2025-05-09 12:29:21 Fri  
**GPT ID:** NA

► Notespace:

### Wilder Blair Munro (AKUSA) said :

**gateway:** 20Wetbrain  
**turn ID:** 761be09a-cba9-4c41-9734-2aeaa3d8f183  
**parent:** f391575e-86b9-4df9-939d-429b1e54a224  
**children:** fa686fcd-ea2b-486c-a3d3-5736ebee14ef  
 ~0-0 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.  
 ~~~~~ ~0-0

Wait...is it possible for me to hardcode a GPT in python? I was under impression that GPTs can only be built through the ChatGPT online builder UI... I am interested in the 'GPT' instead of my own hosted RAG assistant because I do not wish to operate solely through the API, and I wish to be able to share the GPT in the GPT marketplace.

◀ beginning 1 | ◀ prev 3 || this 4 || next 5 ▶ | end 14 ▶▶

Aurora 4omni Munro (ChatGpt) said :

gateway: gpt-4o
turn ID: fa686fcd-ea2b-486c-a3d3-5736ebee14ef
parent: 761be09a-cba9-4c41-9734-2aeaa3d8f183
children: 246a8b38-a015-4b92-b21a-61290744a6e6
 ~~~~~ ~O-O  
 Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.  
 ~~~~~ ~O-O

If you want to build and share a **custom GPT in the OpenAI GPT Builder (ChatGPT UI)** and list it in the **GPT Marketplace**, you must use the **OpenAI ChatGPT online builder UI**. The custom GPTs in this context are distinct from your own custom-hosted models, and they are subject to specific constraints:

Key Points About Custom GPTs in the GPT Builder (OpenAI UI):

1. Web-Based Configuration Only:

- You must use the **ChatGPT online builder UI** to create and configure your custom GPT.
- This includes setting the GPT's instructions, memory, behavior, and knowledge base settings.

2. Knowledge Base Control:

- You can upload files directly in the GPT builder UI as a knowledge base (RAG: Retrieval-Augmented Generation).

- Chunking is managed automatically by OpenAI, but you can influence it using structured text formatting and pre-processing before upload (as we discussed).

3. No Direct Code Execution:

- You cannot directly hardcode a custom GPT using Python and upload it to the GPT Marketplace.
- Your GPT operates within OpenAI's managed environment.

4. Advanced Configuration via GPT Builder UI:

- You can set the behavior, personality, and response style of your GPT through instructions.
- You can also configure its retrieval settings (knowledge base files, API integrations, etc.).

However, If You Want to Simulate a GPT in Python (Separate from GPT Marketplace):

- You can build a **custom RAG (Retrieval-Augmented Generation) assistant** using Python and the OpenAI API.
- This allows you to:
 - Control chunking (as in the Python example I provided).
 - Use vector stores (e.g., Pinecone, Chroma, FAISS).
 - Fine-tune chunking behavior, retrieval logic, and response generation.

Options for Your Goals:

1. To Create a Custom GPT for the Marketplace:

- Use the **ChatGPT builder UI**.
- Upload your knowledge base and structure your file for ideal chunking.
- Configure personality, behavior, and instructions to match your use case.


2. To Build a Custom Python GPT-like Assistant (Not in the Marketplace):

- Use the OpenAI API and Python (with LangChain or similar).
- Control chunking, retrieval, and response behavior manually.
- This approach is flexible but cannot be directly listed in the GPT Marketplace.

Would you like me to show you the **best practices for optimizing your GPT in the GPT Builder UI** so that it performs as close as possible to a custom Python model without API complexity?

I can also show you how to pre-format your knowledge base file to achieve better chunking in the UI.

◀ beginning 1 | ◀ prev 3 || this 4 || next 5 ▶ | end 14 ▶▶



11 / 26

- ### 3. Key Differences Between ChatGPT Pro and API Usage:

4. When to Use Each:

- ## 5. Can You Use Your Pro Account Without Additional API Costs?

- ## 6. Recommendations for Cost Efficiency:

- Would you like me to show you a **full Python example that uses your API key efficiently**, demonstrating how you can run a RAG assistant while minimizing API costs?

This could help you make an informed decision about your use case.

<< markdown file @ line 478 >>



~~~~~



Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.



~~~~~



Conversation 141, Turnpair 6 of 14 ::

@ 2025-05-09 11:56:41 Fri [Δt 0d 00h 05m 57s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]

origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]

created: 2025-05-09 09:47:54 Fri

updated: 2025-05-09 12:29:21 Fri

GPT ID: NA

► Notespace:

Wilder Blair Munro (AKUSA) said :

gateway: 20Wetbrain

turn ID: d5305c0d-a48d-4538-a45d-37014c3d431d

parent: 98c61919-9016-4c5d-a0d5-67003280d523

children: b008c63d-b011-490f-949d-208e2da6fa49

~0-0 ~~~~~

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.

~ ~~~~~ ~0-0

What is the max payload size for a GPT action API return (ie, max tokens to return for injection after tool-call during inference)?

◀◀ beginning 1 | ◀ prev 5 || this 6 || next 7 ▶ | end 14 ▶▶

Aurora 4omni Munro (ChatGpt) said :

gateway: gpt-4o

turn ID: 189c7773-de68-46e6-87b4-ca469b15f085

parent: cbff3d33-e6bd-4a04-ac54-61a9c4bbaa71

children: a7f0dd0d-20eb-491c-8af0-0795f74e30de

~ ~~~~~ ~0-0

Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.

~0-0 ~~~~~

When using GPT Actions (formerly known as plugins) within a custom GPT, the size of the response returned by an action—such as a tool call—is constrained by the model's overall context window. This context window encompasses both the input (prompt) and the output (completion), including any data returned from tools.

🔍 Key Token Limits for GPT Models

| Model | Maximum Context Window (Tokens) | Notes |
|-------------|---------------------------------|--|
| GPT-4.1 | 1,000,000 | Latest model with extended context capabilities citeturn0search23. |
| GPT-4 Turbo | 128,000 | Supports longer contexts citeturn0search22. |

<< markdown file @ line 556 >>

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.  
 ~~~~~

Conversation 141, Turnpair 7 of 14 ::

@ 2025-05-09 11:59:44 Fri [Δt 0d 00h 03m 00s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA

► Notespace:

Wilder Blair Munro (AKUSA) said :

gateway: 20Wetbrain
turn ID: a7f0dd0d-20eb-491c-8af0-0795f74e30de
parent: 189c7773-de68-46e6-87b4-ca469b15f085
children: 90bbe02c-6855-4af3-a63b-ed9c86fa14fd
 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.  
 ~~~~~

I was told by the GPT Builder GPT, that there is an actual size limit to an API return for a GPT action endpoint call.

And, which of the above models is 'gpt 4o'?

◀ beginning 1 | ◀ prev 6 || this 7 || next 8 ▶ | end 14 ▶▶

Aurora 4omni Munro (ChatGpt) said :

gateway: gpt-4o
turn ID: 1a80f885-6f74-42e4-b5c7-f2c93056b16f
parent: 711274a2-45ba-4f48-b159-3c22affd965e
children: 450e413e-d337-4ecf-b209-78f9fa634e64
 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.  
 ~~~~~

You're correct—when using **GPT Actions** (formerly known as plugins) within a custom GPT built through OpenAI's ChatGPT Builder, there are specific constraints on the size of the API responses that can be returned and injected into the model during inference.

GPT Action API Response Size Limits

While OpenAI hasn't published official documentation specifying the exact limits, community observations suggest the following:

- **Maximum Response Size:** Approximately **100 KB** per action response. citeturn0search0
- **Token Considerations:** The total number of tokens (including both input and output) must not exceed the model's context window. For example, if you're using GPT-4o with a 128,000-token context

window, the combined size of the prompt, tool responses, and model output must stay within this limit.

citeturn0search13

Understanding GPT-4o

GPT-4o (the "o" stands for "omni") is a multimodal model released by OpenAI in May 2024. It can process and generate text, images, and audio, making it suitable for a wide range of applications.

- **Context Window:** Supports up to **128,000 tokens**, allowing for extensive input and output within a single interaction. citeturn0search13
 - **Multimodal Capabilities:** Can handle text, images, and audio inputs and outputs, enabling more dynamic and versatile applications. citeturn0search19
-

Best Practices for Managing API Response Sizes

To ensure your GPT Actions function effectively within these constraints:

1. Optimize Data Payloads:

- Ensure that the JSON or data returned by your API is concise and contains only necessary information.

2. Implement Pagination:

- For APIs that return large datasets, implement pagination to retrieve data in manageable chunks.

3. Summarize Large Responses:

- If the API returns extensive information, consider summarizing the data before returning it to the model.

4. Monitor Token Usage:

- Keep track of the number of tokens used in both the prompt and the model's responses to stay within the context window.

5. Handle Errors Gracefully:

- Implement error handling to manage situations where the response size exceeds limits, such as by notifying the user or requesting a more specific query.
-

If you need assistance in designing your GPT Action endpoints to adhere to these constraints or have further questions about GPT-4o's capabilities, feel free to ask!

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

<< markdown file @ line 649 >>

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

Conversation 141, Turnpair 8 of 14 ::

@ 2025-05-09 12:03:04 Fri [4t 0d 00h 03m 18s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA

► Notespace:

Wilder Blair Munro (AKUSA) said :

gateway: 20Wetbrain
turn ID: 450e413e-d337-4ecf-b209-78f9fa634e64
parent: 1a80f885-6f74-42e4-b5c7-f2c93056b16f
children: fbedba82-8b47-4625-b9a9-38d432027cb4
 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.  
 ~~~~~

You gave this:

""""

Model Maximum Context Window (Tokens) Notes

GPT-4.1 1,000,000 Latest model with extended context capabilities .

GPT-4 Turbo 128,000 Supports longer contexts .

GPT-4 (8k) 8,192 Standard context window .

GPT-3.5 Turbo 4,096 Commonly used model with moderate context window .""""

But which one is 4o? 4o is not listed here.

Also, how many characters is 100k of .txt content worth about?

◀ beginning 1 | ◀ prev 7 || this 8 || next 9 ▶ | end 14 ▶▶

Aurora 4omni Munro (ChatGpt) said :

gateway: gpt-4o
turn ID: 430fd4b8-61ad-4445-b3f7-e476ce08d0bf
parent: 3cb1c6f3-f520-41aa-8294-5062cb21df65
children: 08a6d81a-cba4-4bb7-b936-0abe77c71adb
 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.  
 ~~~~~

- **GPT-4o** is not directly listed among the models you mentioned because it is a specialized variant of GPT-4, optimized for multimodal capabilities (text, images, and audio) with a high context window of

128,000 tokens. It is effectively a variant of **GPT-4 Turbo (128k)** but with additional multimodal functionality.

- **100,000 tokens** of text is approximately **400,000 characters**. This is a rough estimate based on an average of **4 characters per token**, which is typical for English text.

Would you like me to show you a practical method to test and validate this for your own text files directly?

◀◀ beginning 1 | ◀ prev 7 || this 8 || next 9 ▶ | end 14 ▶▶

~~~~~

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

~~~~~

<< markdown file @ line 709 >>

~~~~~

Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.

~~~~~

Conversation 141, Turnpair 9 of 14 ::

@ 2025-05-09 12:04:17 Fri [Δt 0d 00h 01m 07s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA

► *Notespace:*

Wilder Blair Munro (AKUSA) said :

[illegible]

No you said "Maximum Response Size: Approximately 100 KB per action response." . How many characters of regular text is this (not tokenized yet)?

⏮ beginning 1 | ◀ prev 8 || this 9 || next 10 ▶ | end 14 ⏭

Aurora 4omni Munro (ChatGpt) said :

gateway: gpt-4o
turn ID: c158fd9e-9bc8-499f-a5db-1b4a58cb9f02
parent: f1087ff1-35b1-4d0e-8454-70c22d496ee0
children: ac7da219-1847-4d6d-b88f-49d4e37dc4db

◆ ~~~~~ ○○~
 Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.
 ◆ ~~~~~ ○○~

A **100 KB** response limit for a GPT Action API call is approximately **102,400 characters** of regular plain text (.txt file).

◀ beginning 1 | ◀ prev 8 || this 9 || next 10 ▶ | end 14 ▶▶

```
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
```

<< markdown file @ line 756 >>

```
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
```

Conversation 141, Turnpair 10 of 14 ::

@ 2025-05-09 12:04:54 Fri [Δt 0d 00h 00m 28s]

```
conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA
```

► Notespace:

Wilder Blair Munro (AKUSA) said :

```
gateway: 20Wetbrain
turn ID: ac7d2a19-1847-4d6d-b88f-49d4e37dc4db
parent: c158fd9e-9bc8-499f-a5db-1b4a58cb9f02
children: e64df57b-a893-40a1-b866-5bf6b6f0a22d
~0-0
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.
~0-0
```

Ah ok. Thank you!

◀ beginning 1 | ◀ prev 9 || this 10 || next 11 ▶ | end 14 ▶▶

Aurora 4omni Munro (ChatGpt) said :

```
gateway: gpt-4o
turn ID: e64df57b-a893-40a1-b866-5bf6b6f0a22d
parent: ac7d2a19-1847-4d6d-b88f-49d4e37dc4db
children: ad79bee6-e8e1-4f2b-a409-53a20f7f6b29
~0-0
Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.
~0-0
```

You're welcome! 😊 Let me know if you have any other questions. Good luck with your GPT development! 🚀

◀ beginning 1 | ◀ prev 9 || this 10 || next 11 ▶ | end 14 ▶▶

```
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
```

<< markdown file @ line 803 >>

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.  
 ~~~~~

Conversation 141, Turnpair 11 of 14 ::

@ 2025-05-09 12:08:41 Fri [4t 0d 00h 03m 45s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA

► Notespace:

Wilder Blair Munro (AKUSA) said :

gateway: 20Wetbrain
turn ID: ad79bee6-e8e1-4f2b-a409-53a20f7f6b29
parent: e64df57b-a893-40a1-b866-5bf6b6f0a22d
children: b777aa8f-f0f8-4270-aec2-21e5ec3a38f8
 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to the children.  
 ~~~~~

Is there some way for me to access the actual final prompt submissions to the LLM? (as in, the GPT instructions, the tool call injected content, the entire 128k, with conversation history, etc.) Looking at the chat export .json, I only really see user input, not the extra tokens.

◀ beginning 1 | ◀ prev 10 || this 11 || next 12 ▶ | end 14 ▶▶

Aurora 4omni Munro (ChatGpt) said :

gateway: gpt-4o
turn ID: b777aa8f-f0f8-4270-aec2-21e5ec3a38f8
parent: ad79bee6-e8e1-4f2b-a409-53a20f7f6b29
children: a862e45d-c36c-402d-9652-86019c67697b
 ~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact in memory of the dead.  
 ~~~~~

Unfortunately, with the **ChatGPT Builder (UI)**, you do not have direct access to the **full final prompt submissions** (including injected tool call content, instructions, and conversation history) that are sent to the LLM. This is because the OpenAI interface is designed to abstract away these details for simplicity and security.

🔍 Why This is the Case:

- OpenAI's ChatGPT Builder UI is a managed environment, which means it handles the assembly of the final prompt (including system instructions, user input, tool call responses, and previous conversation history) behind the scenes.
- The exported **.json** files primarily focus on user input and model output, not the full assembled context window.

- ⚡ **However, There Are Indirect Ways to Access It:**

- If you use the API directly (not the Builder UI), you can log the full prompt you construct.
- This includes the system instructions, user input, tool call responses, and any conversation history you append.

- If you want to see how the final prompt is constructed:
 - Manually replicate the structure in a Python script using the OpenAI API.
 - Log the final assembled prompt before sending it.

- If your GPT is using **Tools or Actions**, you can inspect the API responses of the tool calls before they are injected into the prompt.
- This gives you visibility into the data being injected.

- If you are part of an enterprise or research program, you can request debugging logs for your GPT for specific testing purposes.
- This might give you more visibility into the full context window.

This would allow you to see how your prompt is being assembled without needing direct OpenAI logs.

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.  
 ~~~~~

~~~~~  
 Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.  
 ~~~~~


Conversation 141, Turnpair 14 of 14 ::

@ 2025-05-09 12:29:00 Fri [Δt 0d 00h 01m 02s]

conversation ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
origin ID: 681e3fca-441c-8003-b415-dfae1e6ac25d [conversation 141]
created: 2025-05-09 09:47:54 Fri
updated: 2025-05-09 12:29:21 Fri
GPT ID: NA

► *Notespace:*

Wilder Blair Munro (AKUSA) said :

[illegible]

What if I just want to GPT to echo into chat response what the action returns?

```

    << beginning 1 | < prev 12 || this 13 || next None → | end 14 >>

```

Aurora 4omni Munro (ChatGpt) said :

[illegible]

If you simply want your **GPT to echo the exact text returned by an external tool call (action)** directly in the chat response, you can easily achieve this with a simple method:

💡 How to Make GPT Echo Tool Call Return Text:

1. Ensure Your Tool Call Returns Text Directly:

- When defining your tool or API action, make sure it returns a **plain text response**.
- This text should be exactly what you want the GPT to echo.

2. Modify GPT's Instructions (Prompt Configuration):

- In your GPT Builder (ChatGPT UI):
 - Add a specific instruction such as:

If a tool call returns text, immediately echo that text in your response without any modifications.

- o This forces the GPT to use the returned text as its direct response.

3. Structure Your Tool Call (Action) for Direct Return:

- Ensure your tool call is configured to output text in a clean format, such as:

```
{
  "response": "This is the text returned by the tool."
}
```

- If the tool returns JSON or a more complex object, make sure the GPT is instructed to extract and echo the specific field you need.

4. Example Usage:

- User: "Perform action <<echo_action_response>>."
- Tool (API) Returns: "Hello, this is the response from the tool."
- GPT Responds: "Hello, this is the response from the tool."

Alternative: Explicit Response Template

- If you want even more control, you can pre-define a response template for the GPT, such as:

```
Tool Response: {{tool_response_text}}
```

- In this case, the GPT will always format the response using this template.

Would you like me to show you a version that even auto-appends the tool response with your own customizable prefix (e.g., "Tool says:")?

This is a slightly more sophisticated approach but still very simple to set up.

◀ beginning 1 | ◀ prev 12 || this 13 || next None → | end 14 ▶▶

```
~~~~~
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
~~~~~
```

<< markdown file @ line 1104 >>

```
~~~~~
Artifact released to public domain (Creative Commons Zero). Artifact dedicated to victims of generalized war.
~~~~~
```

EOF